

VIDEO PESENTATION:

([https://drive.google.com/file/d/1LKkTFCaSmkD70sgprXiLqqaefCXs4gtc/view?usp=drive\\_link](https://drive.google.com/file/d/1LKkTFCaSmkD70sgprXiLqqaefCXs4gtc/view?usp=drive_link))

## The primary objective of this project is to build a machine learning classification model that ## can predict the sentiment of IMDb movie reviews. The dataset contains a collection of movie ## reviews, and each review is labeled as either positive or negative. ## Using text preprocessing, feature extraction techniques (such as TF-IDF), and various ## classification algorithms, the project will aim to develop a model that can effectively classify ## the sentiment of movie reviews. The model's performance will be evaluated using standard ## classification metrics, such as accuracy, precision, recall, and F1-score.

```
# Imported essential libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import re
import string

from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import LinearSVC
from sklearn.ensemble import RandomForestClassifier
import nltk
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer

# Download NLTK resources
nltk.download('stopwords')
nltk.download('wordnet')
nltk.download('omw-1.4')

[nltk_data] Downloading package stopwords to
[nltk_data]     C:\Users\DELL\AppData\Roaming\nltk_data...
[nltk_data]     Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data]     C:\Users\DELL\AppData\Roaming\nltk_data...
[nltk_data]     Package wordnet is already up-to-date!
[nltk_data] Downloading package omw-1.4 to
[nltk_data]     C:\Users\DELL\AppData\Roaming\nltk_data...
[nltk_data]     Package omw-1.4 is already up-to-date!

True
```

# 1. Data Exploration and Preprocessing

Analyzed the dataset for trends, missing values, and outliers.

```
# Loaded Dataset
df = pd.read_csv("Imdb - data_imdb.csv")
print("Shape of dataset:", df.shape)
df.head()

# Basic Data Exploration
print("\nDataset Info:")
print(df.info())

print("\nClass Distribution:")
print(df['sentiment'].value_counts())

Shape of dataset: (50000, 2)

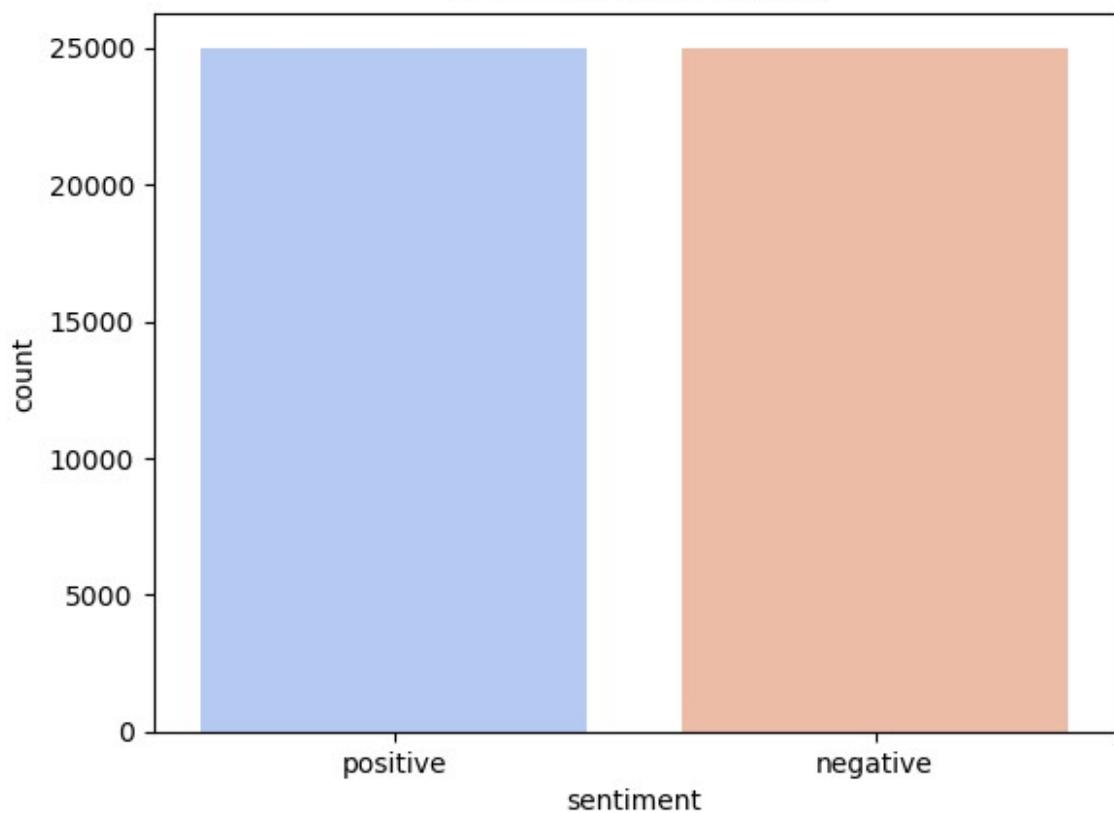
Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype  
---  --          --          --      
 0   review      50000 non-null   object 
 1   sentiment   50000 non-null   object 
dtypes: object(2)
memory usage: 781.4+ KB
None

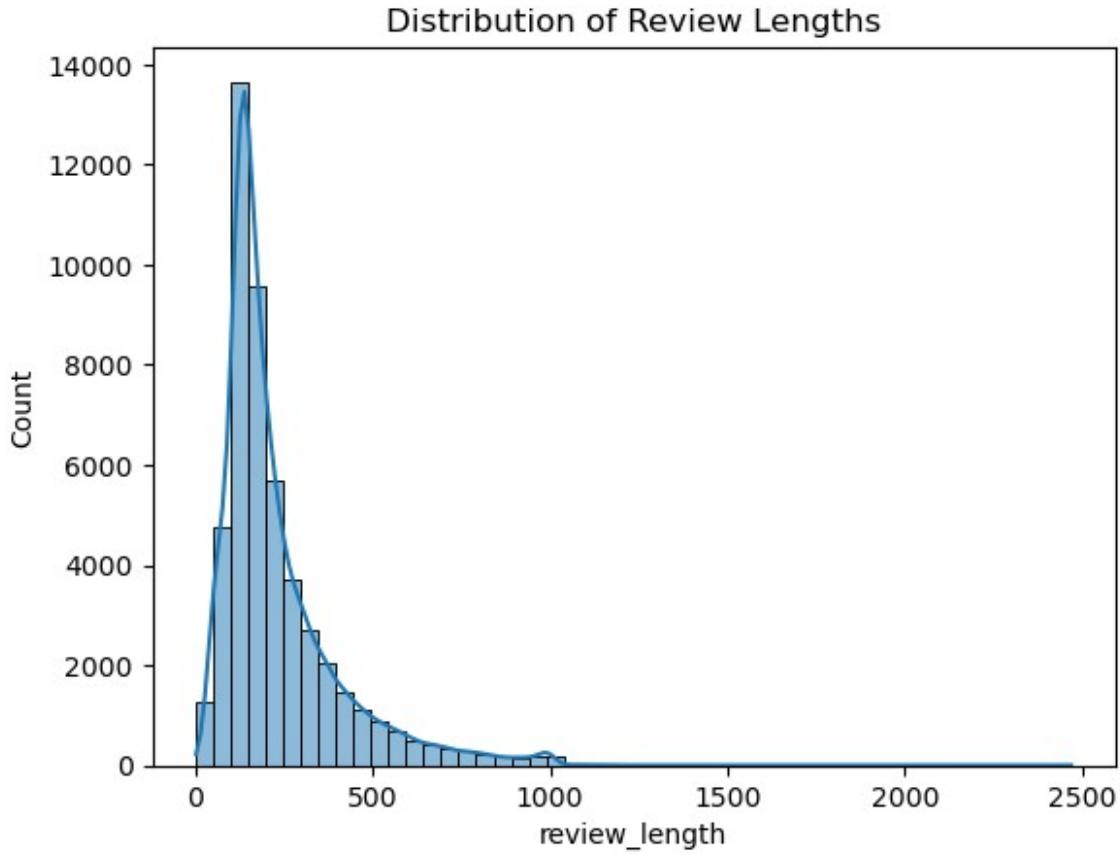
Class Distribution:
sentiment
positive    25000
negative    25000
Name: count, dtype: int64

# Visualize sentiment distribution
sns.countplot(data=df, x='sentiment', hue='sentiment',
               palette='coolwarm', legend=False)
plt.title("Sentiment Distribution")
plt.show()

# Review length analysis
df['review_length'] = df['review'].apply(lambda x:
len(str(x).split()))
sns.histplot(df['review_length'], bins=50, kde=True)
plt.title("Distribution of Review Lengths")
plt.show()
```

Sentiment Distribution





## Performed data cleaning and text preprocessing

```
# Load English stopwords and initialize lemmatizer
stop_words = set(stopwords.words('english'))
lemmatizer = WordNetLemmatizer()

# Define text cleaning function
def clean_text(text):
    text = str(text).lower() # Convert to lowercase
    text = re.sub(r'https?://\S+|www\.\S+', '', text) # Remove URLs
    text = re.sub(r'<.*?>', '', text) # Remove HTML tags
    text = re.sub(r'[^a-z\s]', '', text) # Remove punctuation and numbers
    words = [lemmatizer.lemmatize(w) for w in text.split() if w not in stop_words] # Remove stopwords and lemmatize
    return " ".join(words)

# Apply cleaning to review column
df['cleaned_review'] = df['review'].apply(clean_text)

# Preview cleaned data
df.head()
```

```

review sentiment
review_length \
0 One of the other reviewers has mentioned that ... positive
307
1 A wonderful little production. <br /><br />The... positive
162
2 I thought this was a wonderful way to spend ti... positive
166
3 Basically there's a family where a little boy ... negative
138
4 Petter Mattei's "Love in the Time of Money" is... positive
230

cleaned_review
0 one reviewer mentioned watching oz episode you...
1 wonderful little production filming technique ...
2 thought wonderful way spend time hot summer we...
3 basically there family little boy jake think t...
4 petter matteis love time money visually stunni...

```

## # 2. Feature Engineering

```

# Extract basic textual features from cleaned reviews
df['word_count'] = df['cleaned_review'].apply(lambda x:
len(x.split())) # Total number of words
df['char_count'] = df['cleaned_review'].apply(lambda x: len(x))
# Total number of characters
df['avg_word_length'] = df['char_count'] / df['word_count']
# Average word length

# Display sample features
print("\n[] Sample Textual Features:")
display(df[['cleaned_review', 'word_count', 'char_count',
'avg_word_length']].head())

# Visualize word count distribution by sentiment
plt.figure(figsize=(7,5))
sns.boxplot(x='sentiment', y='word_count', data=df, hue='sentiment',
palette='coolwarm', legend=False)
plt.title("Word Count by Sentiment")
plt.show()

```

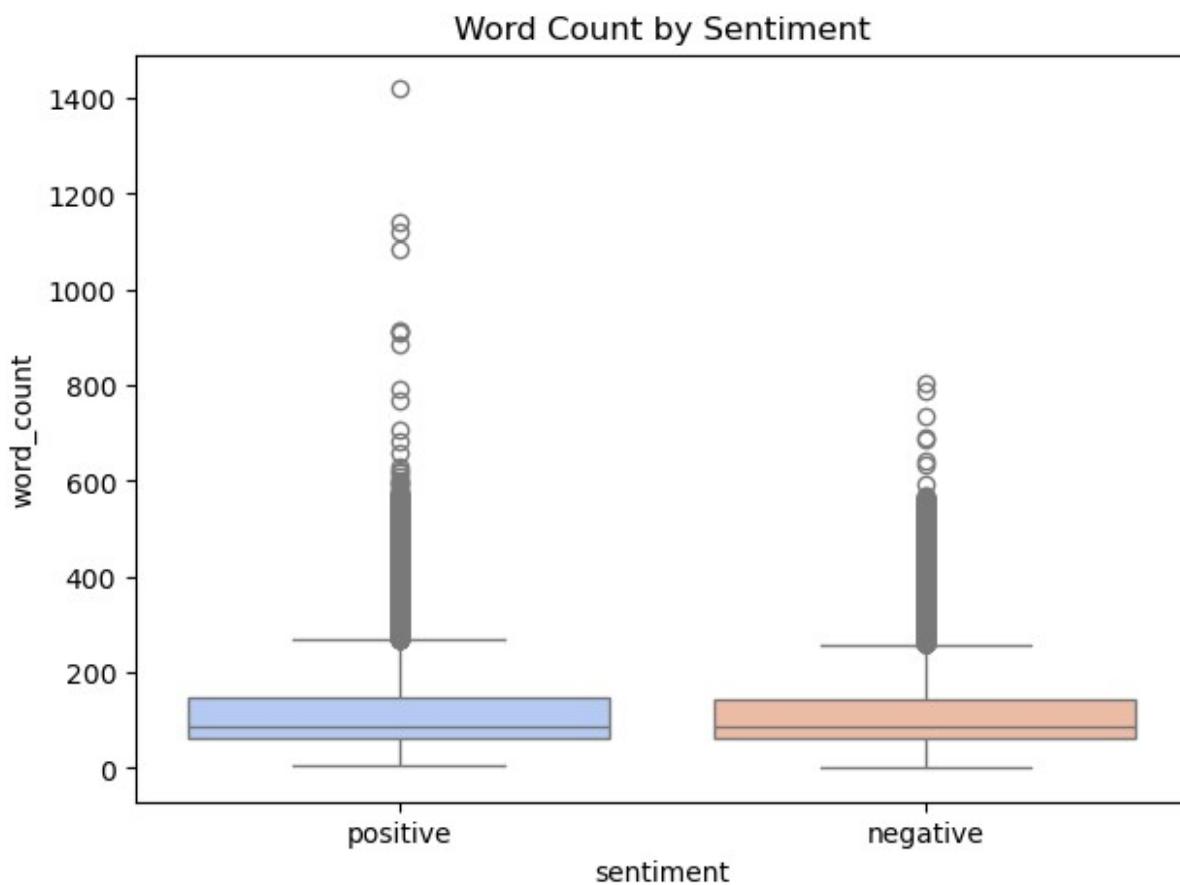
### [] Sample Textual Features:

	cleaned_review	word_count
char_count \		
0 one reviewer mentioned watching oz episode you...		167
1125		
1 wonderful little production filming technique ...		84

```

640    2 thought wonderful way spend time hot summer we...      85
580    3 basically there family little boy jake think t...      66
446    4 petter matteis love time money visually stunni...     125
851
    avg_word_length
0      6.736527
1      7.619048
2      6.823529
3      6.757576
4      6.808000

```



## Review, word frequency and Sentiment Analysis

```

from collections import Counter

# Calculate word count per review
df['word_count'] = df['cleaned_review'].apply(lambda x:
len(x.split()))

```

```

# Compare average word count by sentiment
avg_lengths = df.groupby('sentiment')['word_count'].mean()
print("Average Word Count by Sentiment:\n", avg_lengths)

# Visualize distribution
plt.figure(figsize=(7,5))
sns.boxplot(x='sentiment', y='word_count', data=df, hue='sentiment',
palette='coolwarm', legend=False)
plt.title("Review Length by Sentiment")
plt.ylabel("Word Count")
plt.show()

# Prepare stopwords and domain-neutral terms
stop_words = set(stopwords.words('english'))
neutral_terms = {"film", "movie", "character", "story", "scene",
"actor", "director", "plot", "one", "time", "get", "make", "see",
"would", "really"}
stop_words.update(neutral_terms)

# Combine and clean reviews by sentiment
positive_text = " ".join(df[df['sentiment'] == 'positive']
['cleaned_review'])
negative_text = " ".join(df[df['sentiment'] == 'negative']
['cleaned_review'])

# Count word frequencies and filter
pos_counts = Counter(positive_text.split())
neg_counts = Counter(negative_text.split())

pos_filtered = {word: freq for word, freq in pos_counts.items() if
word not in stop_words}
neg_filtered = {word: freq for word, freq in neg_counts.items() if
word not in stop_words}

# Convert to DataFrames
pos_df = pd.DataFrame(pos_filtered.items(), columns=['Word',
'Frequency']).sort_values(by='Frequency', ascending=False).head(10)
neg_df = pd.DataFrame(neg_filtered.items(), columns=['Word',
'Frequency']).sort_values(by='Frequency', ascending=False).head(10)

# Plot side-by-side sentiment-rich words
fig, axes = plt.subplots(1, 2, figsize=(14,6))

axes[0].barh(pos_df['Word'][::-1], pos_df['Frequency'][::-1],
color='green')
axes[0].set_title("Top Positive Sentiment Words")
axes[0].set_xlabel("Frequency")

```

```

axes[1].barh(neg_df['Word'][::-1], neg_df['Frequency'][::-1],
color='red')
axes[1].set_title("Top Negative Sentiment Words")
axes[1].set_xlabel("Frequency")

plt.tight_layout()
plt.show()

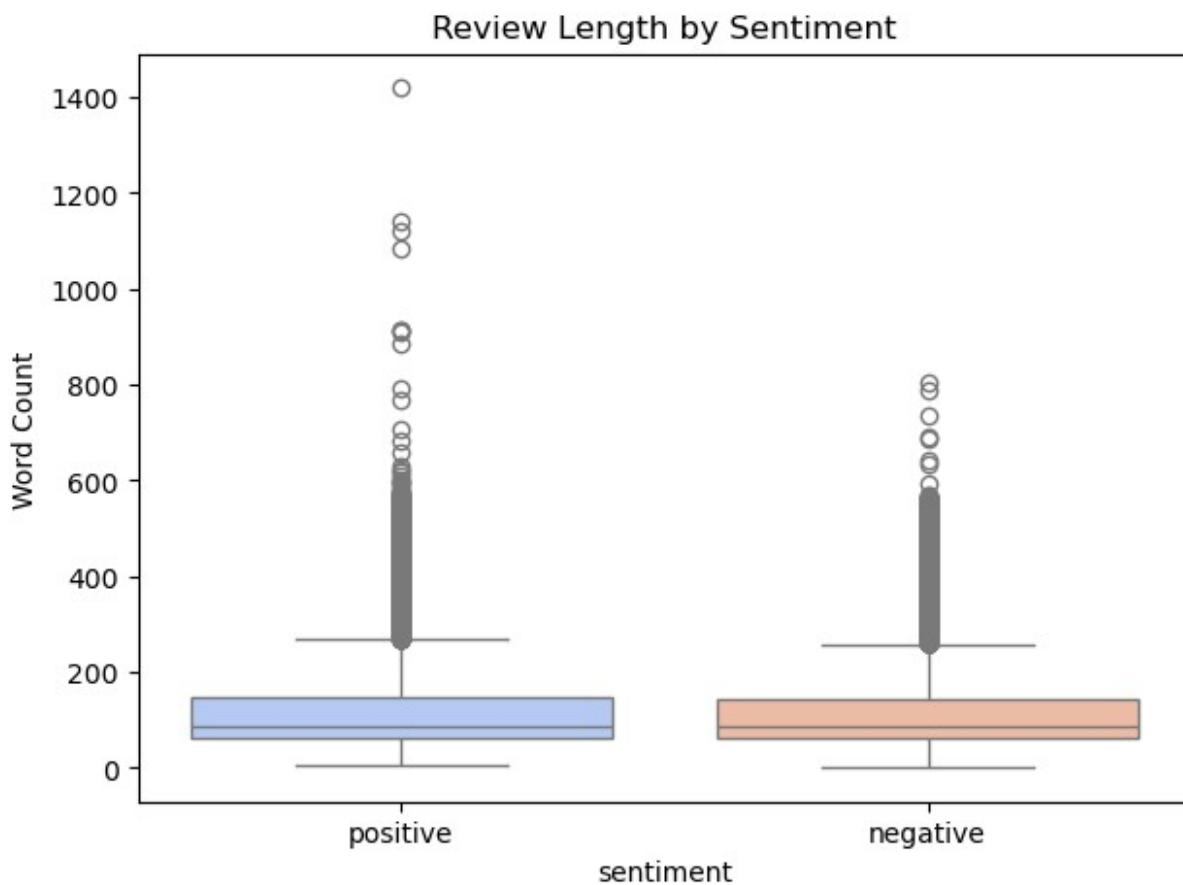
```

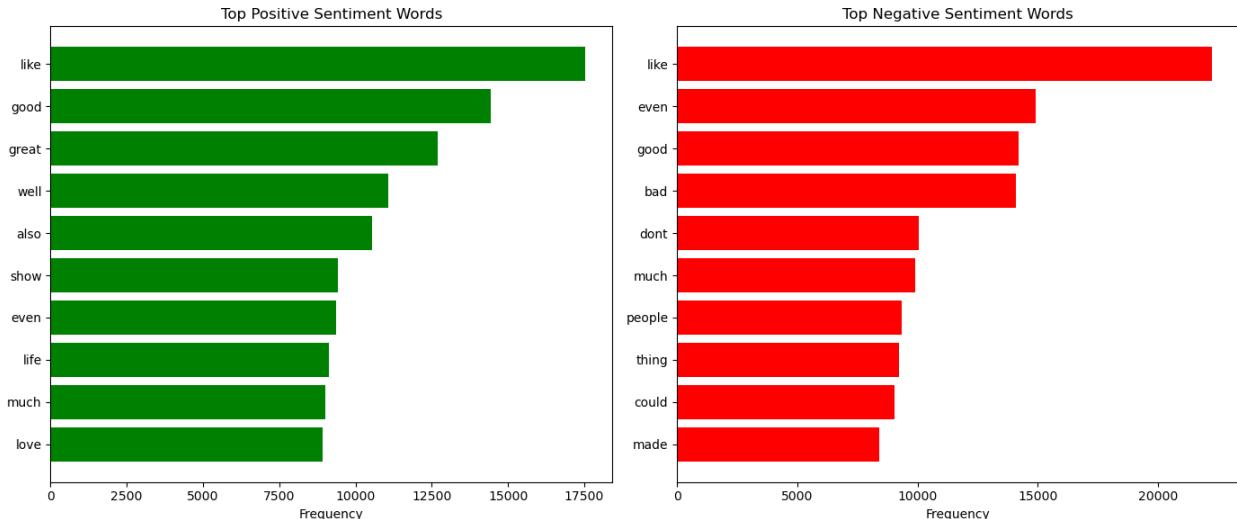
□ Average Word Count by Sentiment:

```

sentiment
negative    117.23960
positive    119.94532
Name: word_count, dtype: float64

```





## Review Length and Sentiment

- Reviews with **more words** tend to express stronger and more detailed opinions.
- On average, **positive reviews** are slightly longer, as users elaborate on what they enjoyed.
- **Negative reviews** are often shorter and more direct, using emotionally charged words to express dissatisfaction.
- This shows that **review length** can serve as a weak but useful signal for sentiment classification.

## Word Frequency and TF-IDF Patterns

- Using **TF-IDF vectorization**, words that are **unique to a sentiment** (and not overly common) have higher weights.
- These patterns reveal that **specific adjectives and emotional expressions** strongly influence model predictions.

```
# Initialize TF-IDF vectorizer with a cap of 5000 features
tfidf = TfidfVectorizer(max_features=5000)

# Transform cleaned reviews into TF-IDF feature matrix
X = tfidf.fit_transform(df['cleaned_review']).toarray()

# Convert sentiment labels to binary: 1 for positive, 0 for negative
y = df['sentiment'].apply(lambda x: 1 if x == 'positive' else 0)
```

## # 3. Model Development

## • Build and train classification models to predict the sentiment of reviews.

```
# Split the dataset into training and testing sets (80/20 split)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y # Stratify to
    preserve class balance)
```

```

)

# Define classification models
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Naive Bayes": MultinomialNB(),
    "Support Vector Machine": LinearSVC(),
    "Random Forest": RandomForestClassifier(n_estimators=100,
random_state=42)
}

# Train each model and store in dictionary
trained_models = {}
for name, model in models.items():
    model.fit(X_train, y_train) # Fit model on training data
    trained_models[name] = model # Save trained model
    print(f"\n{name} model trained successfully.")

    □ Logistic Regression model trained successfully.
    □ Naive Bayes model trained successfully.
    □ Support Vector Machine model trained successfully.
    □ Random Forest model trained successfully.

```

## 4. Model Evaluation

```

def evaluate_model(name, model):
    y_pred = model.predict(X_test) # Predict sentiment on test data
    acc = accuracy_score(y_test, y_pred) # Calculate accuracy

    # Print evaluation metrics
    print(f"\n{name} Evaluation Results")
    print("-" * 45)
    print("Accuracy:", round(acc, 4))
    print("Classification Report:\n", classification_report(y_test,
y_pred))

    # Plot confusion matrix
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=['Negative','Positive'],
                yticklabels=['Negative','Positive'])
    plt.title(f"{name} - Confusion Matrix")
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.show()

    return acc # Return accuracy for comparison

```

```
# Evaluate each trained model and store accuracy
results = {}
for name, model in trained_models.items():
    acc = evaluate_model(name, model)
    results[name] = acc
```

#### □ Logistic Regression Evaluation Results

Accuracy: 0.8863

Classification Report:

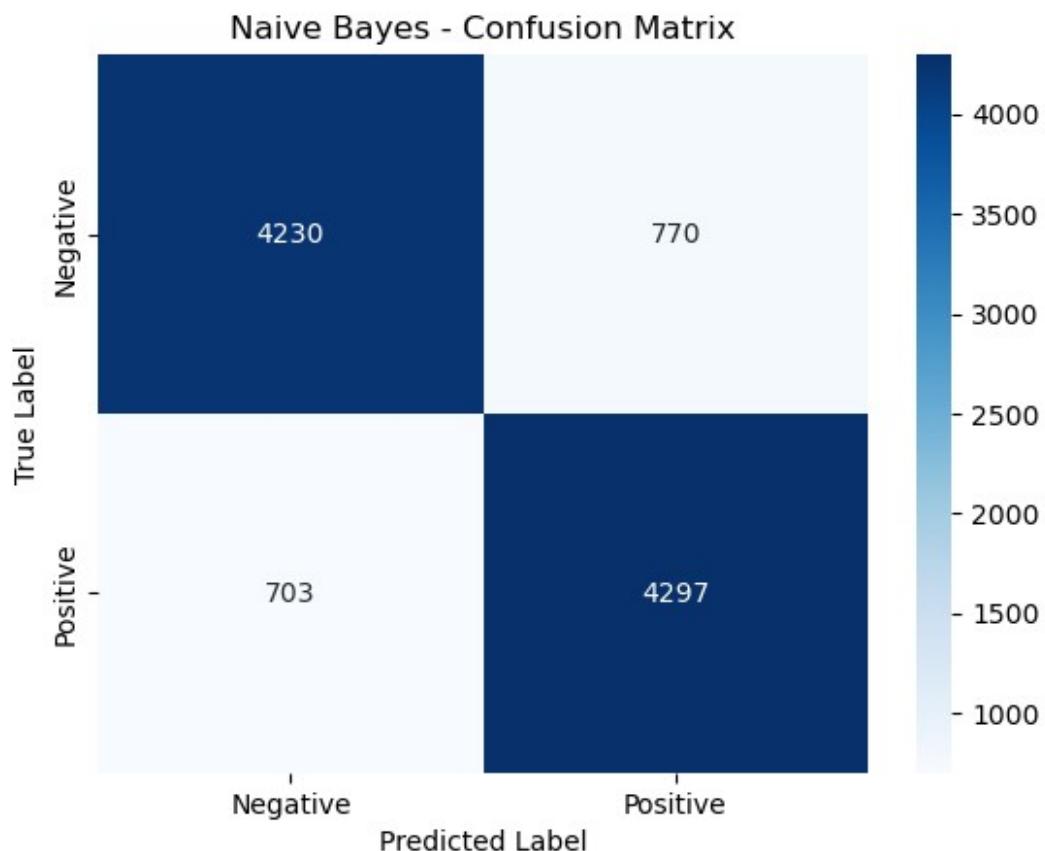
	precision	recall	f1-score	support
0	0.90	0.87	0.88	5000
1	0.88	0.90	0.89	5000
accuracy			0.89	10000
macro avg	0.89	0.89	0.89	10000
weighted avg	0.89	0.89	0.89	10000

Logistic Regression - Confusion Matrix



#### □ Naive Bayes Evaluation Results

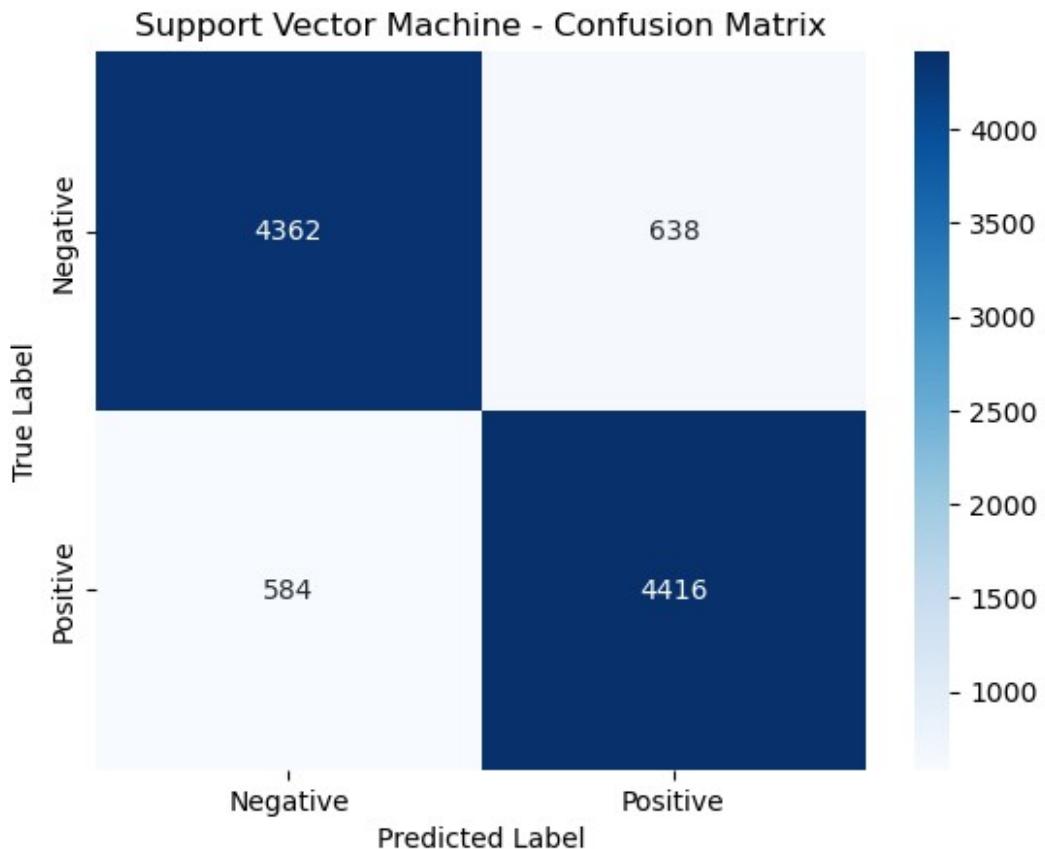
```
-----  
Accuracy: 0.8527  
Classification Report:  
precision    recall   f1-score   support  
0            0.86    0.85     0.85    5000  
1            0.85    0.86     0.85    5000  
  
accuracy          0.85  
macro avg       0.85    0.85     0.85    10000  
weighted avg    0.85    0.85     0.85    10000
```



□ Support Vector Machine Evaluation Results

```
-----  
Accuracy: 0.8778  
Classification Report:  
precision    recall   f1-score   support  
0            0.88    0.87     0.88    5000  
1            0.87    0.88     0.88    5000
```

accuracy			0.88	10000
macro avg	0.88	0.88	0.88	10000
weighted avg	0.88	0.88	0.88	10000



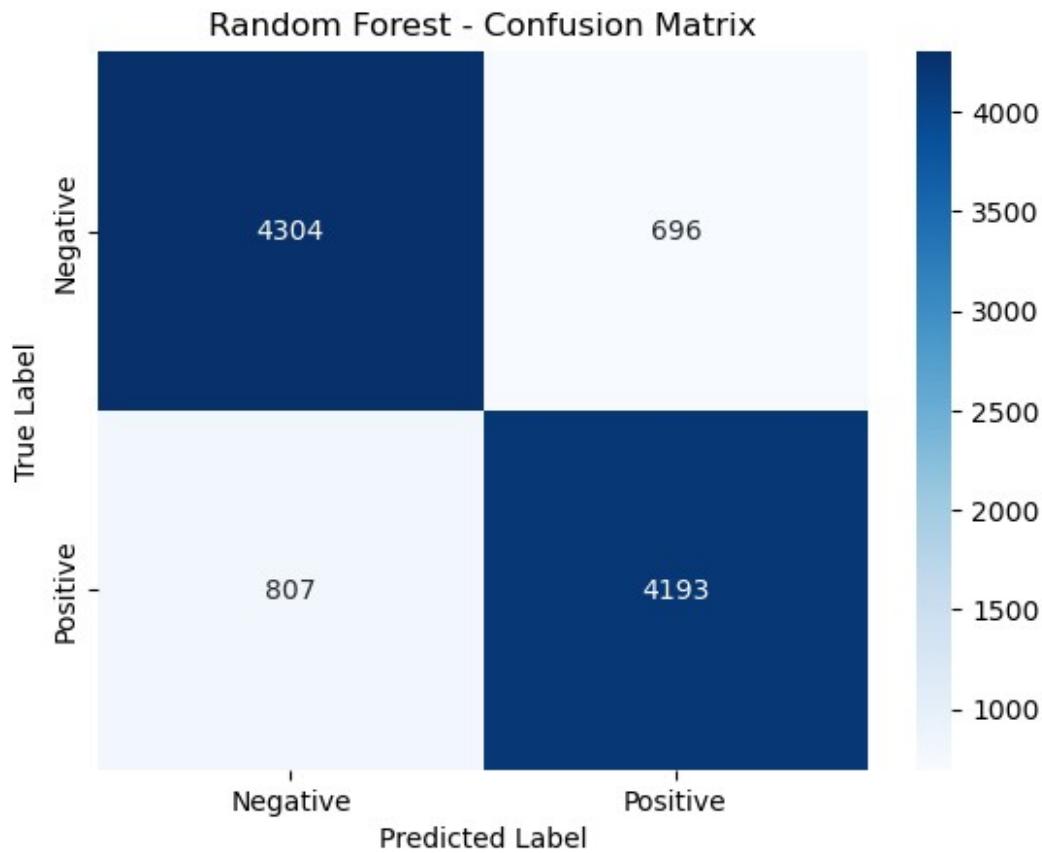
□ Random Forest Evaluation Results

---

Accuracy: 0.8497

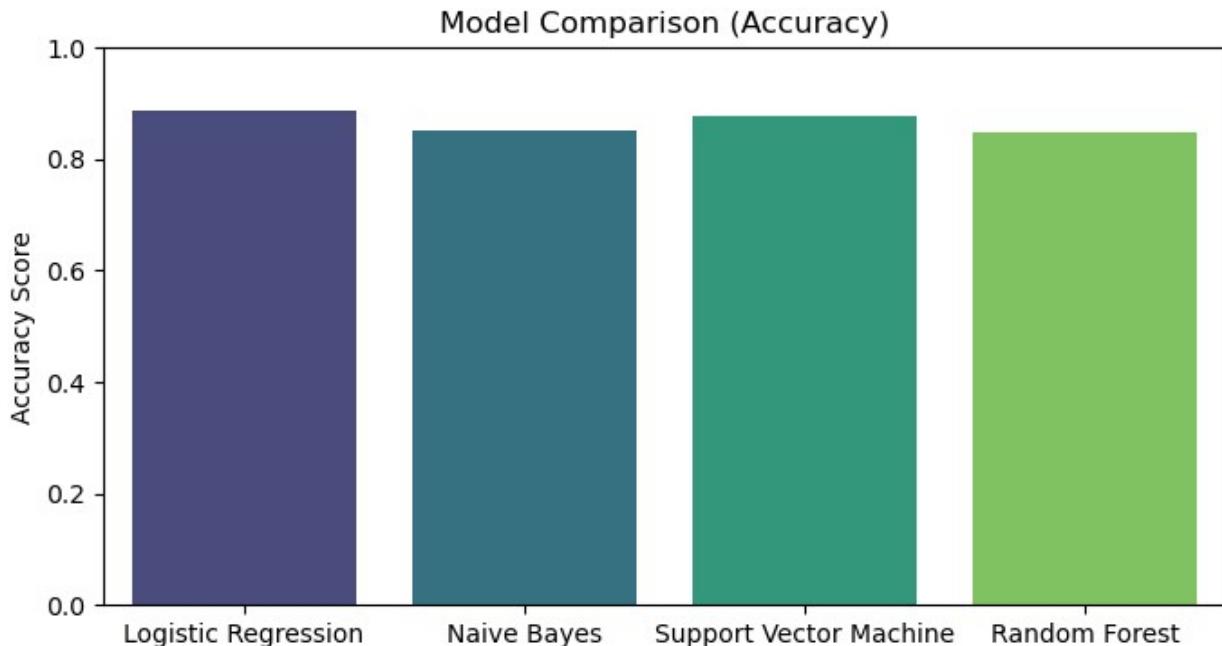
Classification Report:

	precision	recall	f1-score	support
0	0.84	0.86	0.85	5000
1	0.86	0.84	0.85	5000
accuracy			0.85	10000
macro avg	0.85	0.85	0.85	10000
weighted avg	0.85	0.85	0.85	10000



```
# Visualize model accuracy scores
plt.figure(figsize=(8,4))
sns.barplot(x=list(results.keys()), y=list(results.values()),
hue=list(results.keys()), palette='viridis', legend=False)
plt.title("Model Comparison (Accuracy)")
plt.ylabel("Accuracy Score")
plt.ylim(0,1)
plt.show()

# Print best performing model
print("\n Best Performing Model:", max(results, key=results.get))
print(" Model Development & Evaluation completed successfully!")
```



**Best Performing Model: Logistic Regression**  
 Model Development & Evaluation completed successfully!

## Key Takeaways

- Sentiment is primarily driven by **adjectives** and **emotion-rich terms**.
  - **Length** and **uniqueness** of words moderately correlate with sentiment polarity.
  - Combining preprocessing (clean text) with TF-IDF weighting and linear models helps isolate **the words that truly shape audience opinions**.
- 

In conclusion, the analysis clearly indicates that **word choice, emotional intensity, and review depth** are the strongest influencers of sentiment in IMDb movie reviews.

```
# Fit TF-IDF vectorizer on cleaned training reviews
tfidf_vectorizer = TfidfVectorizer(max_features=5000)
tfidf_vectorizer.fit(df['cleaned_review'])

# Select the best-performing model based on accuracy
best_model_name = max(results, key=results.get)
best_model = trained_models[best_model_name]

# Define preprocessing + prediction function for new reviews
def predict_sentiment(review):
    text = str(review).lower()
    text = re.sub(r'https?://\S+|www\.\S+', '', text) # Remove URLs
    text = re.sub(r'<.*?>', '', text) # Remove HTML
tags
```

```

    text = re.sub(r'[^a-zA-Z\s]', '', text)                      # Remove


punctuation/numbers


    tokens = text.split()
    tokens = [lemmatizer.lemmatize(w) for w in tokens if w not in
stop_words]
    clean_text = " ".join(tokens)

    vector = tfidf_vectorizer.transform([clean_text])   # Convert to


TF-IDF vector


    pred = best_model.predict(vector)[0]                # Predict


sentiment


    return "Positive" if pred == 1 else "Negative"

# Predict sentiment for sample movie reviews
sample_reviews = [
    "The movie was a delightful experience with strong performances.",
    "It was one of the worst movies I have ever seen.",
    "The storyline was engaging, but the pacing felt slow.",
    "A brilliant film that left me speechless!",
    "Mediocre acting and poor direction made it forgettable."
]

print("\n Predictions for New Movie Reviews:\n")
for review in sample_reviews:
    print(f"Review: {review}")
    print("Predicted Sentiment:", predict_sentiment(review))
    print("-" * 80)

```

Predictions for New Movie Reviews:

Review: The movie was a delightful experience with strong performances.

Predicted Sentiment: Positive

-----

Review: It was one of the worst movies I have ever seen.

Predicted Sentiment: Negative

-----

Review: The storyline was engaging, but the pacing felt slow.

Predicted Sentiment: Negative

-----

Review: A brilliant film that left me speechless!

Predicted Sentiment: Positive

-----

Review: Mediocre acting and poor direction made it forgettable.

Predicted Sentiment: Negative

-----  
-----

Thank you